



Compute in Vulkan

Jasper Bekkers

OTOY Inc.

October 2017

Intro

- Doing research into Real Time Path Tracing
- Used to do this in CUDA but ran into limitations
 - Single vendor
 - Gpu/Cpu synchronization overhead
 - Couldn't use device enqueue due to runtime linking requirement
 - Limited texturing support
 - No DXT
 - No texturing from Cpu memory
 - No control over when dispatches to the Gpu happen, so lots of bubbles
- CUDA had major advantages too
 - Full and stable C++ compiler
 - Ease of use
 - Stable API and kernel binary format



What do our use cases look like

- Lots of multi-gpu with tight communication between devices
- Few but long indirect dispatches
- Huge assets so
 - Lots of migrating data, mostly in buffer objects
 - PCIe performance still not good enough
 - Lots of textures too
 - Mostly static on the Gpu
- 100% desktop / server focussed
- Not using graphics pipe at all
 - Do use swapchains & presents but that's it
 - Write directly to swapchain from compute (except on Intel, where we can't)



What does our API look like

- Exposed DescriptorSet as first class citizen within the API
 - Different systems create their own 'global' DescriptorSet
 - KernelArgs object that contains all bindings for a kernel
 - No more SetBuffer calls on the command buffer
- Exposed Heaps and Memory objects
 - Allows for more complex memory usage within the API
 - MemoryTypeIndex made things more difficult
- Exposed command buffers and synchronization
 - Semaphores and Fence's used liberally throughout engine
 - GpuJobSystem is replacing most of the explicit synchronization



Job system

- Automatic synchronization and reordering of compute jobs and copies
 - Supports async compute
 - Supports async copies
- Simple as possible API for setting up jobs
 - Not much different than our regular API wrapped in a lambda
- Resource tracking through opaque memory handles
 - Most resources involved are transient
 - Memory can be reused after jobs end
 - No tiled resources tricks



Vulkan issues

- Bindless works but has usability issues
 - Can't update bindings table sparsely without validation layer going complaining
 - Still no NonUniformResourceIndex equivalent
- Headless servers
- Validation layer is slow and has huge memory overhead
 - Frame that normally takes 800us now takes 22ms
- Missing equivalent to D3D12's SetStablePowerState
 - Hacked around it by calling into D3D12



Vulkan issues

- MemoryTypeIndex not transparent enough
 - It's not clear what buffer or texture will belong to which memory type
 - Makes it more tricky to write allocators against
- Can't share Cpu side buffers between devices (of different vendors)
 - Nor with devices of the same vendor
 - Linear buffer only would be ok
 - VK_KHR_device_group not what we want
 - Don't want masking in command buffer; need to have one Cpu thread per device
 - Nvidia only
 - Adds up when you need to have one staging copy per Gpu in RAM



Vulkan looking ahead

- Want our entire path tracer driven from Gpu
 - Need to do a lot of decision making now that could be kept on the Gpu with lower latency
- Would love to see a `vkCmdCopy*Indirect`
 - Gpu for us already decides what needs to be rendered next, but right now we need a Cpu roundtrip to get the data back in place
 - On the copy queue
 - Only need control over offsets and size
- Would love to see loops in the command buffer
 - Mantle had this as `grCmdWhile` / `grCmdEndWhile` and it's an extremely useful building block
 - This would get rid of most of our Cpu roundtrips
- Would love to see copies between physical devices
 - Cross vendor would be ideal, but even within the same vendor would be great progress



Shading language looking ahead

- Want to start linking SPIR-V
 - Our material system would be our primary customer for this
 - Want to work closely with Pierre to make this happen
- Need to get rid of std430 and std140! Tried to add tight packing to glslang ourselves but need vendor support.
 - Have lots of code that we share with C++ so packing and alignment rules should be similar (at least to VS2017) or enforceable through #pragma's
 - Want to #include shared headers in C++ and GLSL, these mostly contain shared data structures
- Would love more explicit control over loads and stores
 - Need to use vec4, vec2 etc as hints of what stores we want on some vendors
 - I think type is the wrong place to indicate stores
 - Load/store coalescing doesn't happen at all on some vendors



Shading language looking ahead

- Would love to use workgroup sizes & count to enforce spilling and resource allocation
 - Gives easiers control over parallelism
 - CUDA has this; it would influence register allocation, shared memory size etc
 - Was amazing for performance stability (and performance tweaking)
- Would like to have dynamic indexing of uniforms in Vulkan
- Would love to be able to pass parameters into the backend compiler (-ffast-math, unrolling, register counts etc)
- Would love a mechanism for feedback from backend compiler (perf warnings, invalid SPIR-V, disassembly etc)



Shading language problems

- Would love to index into different buffers to get around the 2-4GB buffer limit size
 - Right now we need to store data together (or separate) that we don't really want to
- Complex data structures often break the compiler / build
 - Sometimes break with dynamic loops
 - Had to manually unroll some control flow to get around issues
- Would like references to memory
 - `vec3& ref = someBuffer[i].someMember[j].someOtherMember;`
 - `ref.x = 1;`
 - `// skip y`
 - `ref.z = 3;`



Ecosystem issues

- We've had some difficulty getting feedback after issues have been reported to the VAP
- LunarG don't seem to work on the most useful features first
 - Examples include 6+ month old issue on having names in the validation layer
 - Not validating extremely basic physical device limits
- Would love to contribute driver regressions / repro cases to CTS but its code-base is a huge mess
 - Ideally would want to just build my repro cases in CTS first
 - Should be small and easy to work with and upstream



Questions?



Thanks

Baldur Karlsson

Matthäus G. Chajdas

Neil Henning

Jeroen van Schijndel

Juul Joosten

