

# Implementing a linker for SPIR-V binaries

Pierre Moreau

13th October 2017

# Some background on the project

**Goal** adding OpenCL support to Nouveau (open-source driver for NVIDIA cards, on Linux)

**How** using SPIR-V as a representation of OpenCL to be consumed by Nouveau

**Time frame**

- ▶ Started on consuming SPIR-V in July 2015;
- ▶ Started on the SPIR-V linker in January 2017.

# Some background on the library: SPIRV-Tools

Available at

<https://github.com/KhronosGroup/SPIRV-Tools>. Improved daily by employees from Google and LunarG, among others.

Tools available:

- ▶ Assembler, from textual to binary representation;
- ▶ Disassembler, from binary to textual;
- ▶ Validator (WIP);
- ▶ Optimiser (WIP);
- ▶ Linker (WIP);
- ▶ Some utilities.

# Introducing SPIR-V

# SPIR-V: An introduction

## Some history

Intermediate Language from the Khronos Group, represented as a stream of 32-bit words, used as input for various APIs:

**OpenCL** Kernels in OpenCL  $\geq 2.1$  should be created from SPIR-V modules.

**OpenGL** Available through the extension ARB\_gl\_spirv, core since 4.6.

**Vulkan** Shaders should be created from SPIR-V modules.

Possible to compile from OpenCL, GLSL, HLSL, among others, to SPIR-V.

SPIR-V instruction: `insn_descr operand1 operand2 ...`  
where `insn_descr` contains the instruction opcode and the number of words it is made of.

# SPIR-V: An introduction

## Format, Part 1

The binary is made of a header, followed by all the instructions. Those instructions are grouped in different categories:

1. Capabilities required by the binary;
2. Declaration of use of extensions;
3. Import of extended set of instructions;
4. Declaration of the memory model used;
5. Declaration of all entry points;
6. Declaration of all execution modes;
7. Debug information;
8. Annotations;
9. Declaration of types, constants and global variables;
10. Declaration of function prototypes;
11. Definition of functions.

# SPIR-V: An introduction

## Format, Part 2

```
; SPIR-V
; Version: 1.0
; Generator: Khronos LLVM/SPIR-V Translator; 14
; Bound: 10
; Schema: 0

OpCapability Addresses
OpCapability Kernel
%1 = OpExtInstImport "OpenCL.std"
OpMemoryModel Physical64 OpenCL
OpEntryPoint Kernel %6 "test"
OpSource OpenCL_C 102000
OpName %res "res"
%uint = OpTypeInt 32 0
%uint_42 = OpConstant %uint 42
%void = OpTypeVoid
%_ptr_CrossWorkgroup_uint = OpTypePointer CrossWorkgroup %uint
%5 = OpTypeFunction %void %_ptr_CrossWorkgroup_uint
%6 = OpFunction %void None %5
%res = OpFunctionParameter %_ptr_CrossWorkgroup_uint
%8 = OpLabel
OpStore %res %uint_42 Aligned 4
OpReturn
OpFunctionEnd
```

Figure: The SPIR-V of a simple OpenCL kernel which writes the value 42 at some location in global memory.

## About that linker



# Linking: Overview

The different linking steps are:

1. Convert all binaries to the same endianness (using the one from the host);
2. Generate the header of the resulting module;
3. Shift the IDs in each module;
4. Combine each category from the different modules;
5. Extract the import and export pairs, and check their signature;
6. Remove duplicate information;
7. Remove no longer needed linkage information;
8. Remove imported definitions and replace their usage by the exported declarations;
9. Compact the IDs.

# Linking: Generate the new header

## Header format:

- ▶ Magic word;
- ▶ SPIR-V version number;
- ▶ Generator's magic number;
- ▶ ID bound ( $0 < id < bound$ );
- ▶ Reserved (should be 0);

## Merging issues:

**Version** Currently  $\max(\$version_1, \$version_2, \dots)$ , but it could fail if deprecation is introduced in SPIR-V.

**IDs** Need to shift IDs of each module to prevent them from overlapping each other.

# Linking: Matching imports and exports

## How to identify imported and exported symbols?

- ▶ Use the LinkageAttribute annotation: gives the type of linkage, and a name;
- ▶ An import/export pair shares the same name.

## There is a match if

- ▶ they have the same decorations;
- ▶ they have the same types/signature; types need to have the same decorations to match;
- ▶ their arguments have the same decorations (maybe?); only valid for functions.

# Linking: Cleaning up

The following is done automatically:

- ▶ Remove duplicates:
  - ▶ Capabilities;
  - ▶ Import of extended set of instructions;
  - ▶ Annotations: an import and the corresponding export have the same ones;
  - ▶ Types: idem.
- ▶ Remove declaration of imported global variables and functions;
- ▶ Change code to use IDs of exported symbols, rather than imported ones;
- ▶ Compact IDs: the previous operation will leave some unused IDs;

You might want to manually run some optimisations on the resulting module, such as function inlining, dead code removal, etc.

To conclude

# Conclusion: Current status

## Done:

- ▶ Combine several SPIR-V modules together;
- ▶ (Implement a very basic annotation manager);
- ▶ (Remove various duplicates due to combining modules).

## To do:

- ▶ (Remove duplicate extensions and names);
- ▶ (Remove dead instructions);
- ▶ Support OpenCL maths optimisation flags;
- ▶ Optimise the code (quadratic behaviours);
- ▶ Keep the DefUseManager up-to-date;
- ▶ Other code clean-ups.

# Why is an annotation manager helpful?

## Some facts

- ▶ Annotations can be applied to a single target (type, variable, function, annotation group).
- ▶ Annotation groups can be applied to multiple targets at once.

Groups make it complicated to interact with decorations: you could have T1 with decorations A, B and  $G1\{C, D\}$ , and T2 with A,  $G2\{B\}$ , C, D.

# Conclusion: Current status

## Done:

- ▶ Combine several SPIR-V modules together;
- ▶ (Implement a very basic annotation manager);
- ▶ (Remove various duplicates due to combining modules).

## To do:

- ▶ (Remove duplicate extensions and names);
- ▶ (Remove dead instructions);
- ▶ Support OpenCL maths optimisation flags;
- ▶ Optimise the code (quadratic behaviours);
- ▶ Keep the DefUseManager up-to-date;
- ▶ Other code clean-ups.