

# Introduction to Vulkan Layers

Baldur Karlsson

A dark blue diagonal gradient bar that starts from the bottom left and extends towards the top right, covering the lower half of the slide.

# Vulkan Layers

- Allows additional functionality to be added above what a driver implements.

# Vulkan Layers

- Allows additional functionality to be added above what a driver implements.
- Layers can add, modify or examine all aspects of the API. Layers can even implement extensions that are not supported by the driver.

# Vulkan Layers

- Allows additional functionality to be added above what a driver implements.
- Layers can add, modify or examine all aspects of the API. Layers can even implement extensions that are not supported by the driver.
- Pay only for what you use, deactivated layers have zero cost.

# Vulkan Layers

- Allows additional functionality to be added above what a driver implements.
- Layers can add, modify or examine all aspects of the API. Layers can even implement extensions that are not supported by the driver.
- Pay only for what you use, deactivated layers have zero cost.
- Primarily for during development, but some layers are used with released products like steam overlay.

# Dispatchable Objects

- The loader is the the vulkan dispatch library. vulkan-1.dll or libvulkan1.so.

# Dispatchable Objects

- The loader is the the vulkan dispatch library. vulkan-1.dll or libvulkan1.so.
- It doesn't implement anything itself, it is responsible for enumerating which drivers are installed and calling into them.

# Dispatchable Objects

- The loader is the the vulkan dispatch library. vulkan-1.dll or libvulkan1.so.
- It doesn't implement anything itself, it is responsible for enumerating which drivers are installed and calling into them.
- A VkInstance can have several VkPhysicalDevices - each with a different driver.



# Dispatchable Objects

- The loader is the the vulkan dispatch library. vulkan-1.dll or libvulkan1.so.
- It doesn't implement anything itself, it is responsible for enumerating which drivers are installed and calling into them.
- A VkInstance can have several VkPhysicalDevices - each with a different driver.
- When you call a function like vkCmdDraw() - which driver does it go to?

# Dispatchable Objects

- The loader is the the vulkan dispatch library. vulkan-1.dll or libvulkan1.so.
- It doesn't implement anything itself, it is responsible for enumerating which drivers are installed and calling into them.
- A VkInstance can have several VkPhysicalDevices - each with a different driver.
- When you call a function like vkCmdDraw() - which driver does it go to?
- Answer - it uses whichever driver the objects are set to **dispatch** to.

# Dispatchable Objects

- The loader is the the vulkan dispatch library. vulkan-1.dll or libvulkan1.so.
- It doesn't implement anything itself, it is responsible for enumerating which drivers are installed and calling into them.
- A VkInstance can have several VkPhysicalDevices - each with a different driver.
- When you call a function like vkCmdDraw() - which driver does it go to?
- Answer - it uses whichever driver the objects are set to **dispatch** to.
- Vulkan objects are **dispatchable** or **non-dispatchable**. Dispatchable objects are:
  - VkInstance, VkPhysicalDevice, VkDevice, VkQueue, VkCommandBuffer

# Dispatchable Objects

- Dispatchable objects are always pointer-sized.

# Dispatchable Objects

- Dispatchable objects are always pointer-sized.
- Look again at the list. It's no coincidence that **every** vulkan function call takes a dispatchable object as its first parameter:
  - VkInstance, VkPhysicalDevice, VkDevice, VkQueue, VkCommandBuffer

# Dispatchable Objects

- Dispatchable objects are always pointer-sized.
- Look again at the list. It's no coincidence that **every** vulkan function call takes a dispatchable object as its first parameter:
  - VkInstance, VkPhysicalDevice, VkDevice, VkQueue, VkCommandBuffer
- Each object is created from a parent - command buffers created from a device, so at each point we know which driver this object is associated with.

# Dispatch tables

- Dispatchable objects are pointer-sized and the memory they point to is mostly driver-defined, as with non-dispatchable objects.

# Dispatch tables

- Dispatchable objects are pointer-sized and the memory they point to is mostly driver-defined, as with non-dispatchable objects.
- The loader-layer-driver interface defines that the pointed-to memory contains at its beginning a pointer to a **dispatch table**.



# Dispatch tables

- Dispatchable objects are pointer-sized and the memory they point to is mostly driver-defined, as with non-dispatchable objects.
- The loader-layer-driver interface defines that the pointed-to memory contains at its beginning a pointer to a **dispatch table**.
- A dispatch table is a big list of function pointers, with one for every function call in the vulkan API.

# Dispatch tables

- Dispatchable objects are pointer-sized and the memory they point to is mostly driver-defined, as with non-dispatchable objects.
- The loader-layer-driver interface defines that the pointed-to memory contains at its beginning a pointer to a **dispatch table**.
- A dispatch table is a big list of function pointers, with one for every function call in the vulkan API.
- Inside vkCmdDraw in the loader, it looks up this table to find which function pointer to call - into the right driver for this object.

# Dispatch tables

- Dispatchable objects are pointer-sized and the memory they point to is mostly driver-defined, as with non-dispatchable objects.
- The loader-layer-driver interface defines that the pointed-to memory contains at its beginning a pointer to a **dispatch table**.
- A dispatch table is a big list of function pointers, with one for every function call in the vulkan API.
- Inside vkCmdDraw in the loader, it looks up this table to find which function pointer to call - into the right driver for this object.
- The loader automatically populates these tables behind the scenes. There are two different tables - one for VkInstance objects (driver independent), one for VkDevice objects (driver dependent).

# Dispatch chains

- This solves the problem of how to get `vkCmdDraw()` to the right driver - but how do layers fit in?

# Dispatch chains

- This solves the problem of how to get `vkCmdDraw()` to the right driver - but how do layers fit in?
- There's no reason that this table has to go directly to the driver. It could call a function in a layer library first which then knows to go to the driver:
  - Loader -> Layer -> Driver

# Dispatch chains

- This solves the problem of how to get `vkCmdDraw()` to the right driver - but how do layers fit in?
- There's no reason that this table has to go directly to the driver. It could call a function in a layer library first which then knows to go to the driver:
  - Loader -> Layer -> Driver
- Instead of the table calling directly to the driver, you can call one or more layers in a chain of calls before ending up at the driver:
  - Loader -> Layer A -> Layer B -> Layer C -> Driver

# Dispatch chains

- This solves the problem of how to get `vkCmdDraw()` to the right driver - but how do layers fit in?
- There's no reason that this table has to go directly to the driver. It could call a function in a layer library first which then knows to go to the driver:
  - Loader -> Layer -> Driver
- Instead of the table calling directly to the driver, you can call one or more layers in a chain of calls before ending up at the driver:
  - Loader -> Layer A -> Layer B -> Layer C -> Driver
- For every function in the vulkan API, we then have a chain like this - the application calls into the loader, which calls along all the layers, then finally to the driver.

# Implementing a layer

- A simple layer is quite easy to implement.



# Implementing a layer

- A simple layer is quite easy to implement.
- First you need a JSON manifest registered in a specific location to tell the loader that your layer is available, and where the library file is.

# Implementing a layer

- A simple layer is quite easy to implement.
- First you need a JSON manifest registered in a specific location to tell the loader that your layer is available, and where the library file is.
- Then your layer must implement a few entry points, the most important of which are `vkGetInstanceProcAddr` and `vkGetDeviceProcAddr`.

# Implementing a layer

- A simple layer is quite easy to implement.
- First you need a JSON manifest registered in a specific location to tell the loader that your layer is available, and where the library file is.
- Then your layer must implement a few entry points, the most important of which are `vkGetInstanceProcAddr` and `vkGetDeviceProcAddr`.
- The loader uses this to find out which functions your layer implements using these functions as normal.

# Implementing a layer

- A simple layer is quite easy to implement.
- First you need a JSON manifest registered in a specific location to tell the loader that your layer is available, and where the library file is.
- Then your layer must implement a few entry points, the most important of which are `vkGetInstanceProcAddr` and `vkGetDeviceProcAddr`.
- The loader uses this to find out which functions your layer implements using these functions as normal.
- Inside `vkCreateInstance` or `vkCreateDevice`, your layer can find the `vkGetInstanceProcAddr` and `vkGetDeviceProcAddr` for the next element in the chain (layer or driver) through the `pNext` extension chain.

# Storing Layer Data

- Each link in the dispatch chain has a unique dispatch table, containing the functions of the next link in the chain where it needs to call.

# Storing Layer Data

- Each link in the dispatch chain has a unique dispatch table, containing the functions of the next link in the chain where it needs to call.
- But the dispatch table stored inside the dispatchable object is reserved for the loader, so that it can start the chain. Layers cannot store their own table here.

# Storing Layer Data

- Each link in the dispatch chain has a unique dispatch table, containing the functions of the next link in the chain where it needs to call.
- But the dispatch table stored inside the dispatchable object is reserved for the loader, so that it can start the chain. Layers cannot store their own table here.
- There are two primary ways to store layer-specific data:
  1. Map Lookups
  2. Object Wrapping

# Storing Layer Data – Map Lookups

- The loader's device or instance table is unique for each VkInstance chain or VkDevice chain (newly malloc'd each time).



# Storing Layer Data – Map Lookups

- The loader's device or instance table is unique for each VkInstance chain or VkDevice chain (newly malloc'd each time).
- You can use a `std::map<>` or similar to look up your own dispatch table using the loader's table as a key.

# Storing Layer Data – Map Lookups

- The loader's device or instance table is unique for each VkInstance chain or VkDevice chain (newly malloc'd each time).
- You can use a `std::map<>` or similar to look up your own dispatch table using the loader's table as a key.
- Similarly you can use the object handle itself as a key, and use a `std::map<>` to store object-specific data you may need to look up.

# Storing Layer Data – Object Wrapping

- Layers can ‘wrap’ all of the vulkan objects.

# Storing Layer Data – Object Wrapping

- Layers can ‘wrap’ all of the vulkan objects.
- When a new object - e.g. `VkCommandBuffer` - is created. The layer allocates new memory, stores the loader’s dispatch table at the start, the real object handle next, and after that as much data as it wants

# Storing Layer Data – Object Wrapping

- Layers can ‘wrap’ all of the vulkan objects.
- When a new object - e.g. `VkCommandBuffer` - is created. The layer allocates new memory, stores the loader’s dispatch table at the start, the real object handle next, and after that as much data as it wants:
  - An ID
  - Its dispatch table
  - A pointer to cold data, etc

# Storing Layer Data – Object Wrapping

- Layers can ‘wrap’ all of the vulkan objects.
- When a new object - e.g. `VkCommandBuffer` - is created. The layer allocates new memory, stores the loader’s dispatch table at the start, the real object handle next, and after that as much data as it wants:
  - An ID
  - Its dispatch table
  - A pointer to cold data, etc
- Any time an object is passed down through the API, the layer ‘unwraps’ the object by looking up the real handle and replacing it for further down the chain.

# Storing Layer Data – Summary

## Map Lookups

- Simple to implement.

# Storing Layer Data – Summary

## Map Lookups

- Simple to implement.
- Non-modifying - only requires lookups.



# Storing Layer Data – Summary

## Map Lookups

- Simple to implement.
- Non-modifying - only requires lookups.
- Thread-unsafe - requires read-write lock for object creation & destruction.

# Storing Layer Data – Summary

## Map Lookups

- Simple to implement.
- Non-modifying - only requires lookups.
- Thread-unsafe - requires read-write lock for object creation & destruction.
- Per-object map can have problems with duplicate keys (GOOGLE\_unique\_objects layer).

# Storing Layer Data – Summary

## Map Lookups

- Simple to implement.
- Non-modifying - only requires lookups.
- Thread-unsafe - requires read-write lock for object creation & destruction.
- Per-object map can have problems with duplicate keys (GOOGLE\_unique\_objects layer).

## Object Wrapping

- Complex to implement.

# Storing Layer Data – Summary

## Map Lookups

- Simple to implement.
- Non-modifying - only requires lookups.
- Thread-unsafe - requires read-write lock for object creation & destruction.
- Per-object map can have problems with duplicate keys (GOOGLE\_unique\_objects layer).

## Object Wrapping

- Complex to implement.
- Modifying - layer must intercept all entry points to unwrap.

# Storing Layer Data – Summary

## Map Lookups

- Simple to implement.
- Non-modifying - only requires lookups.
- Thread-unsafe - requires read-write lock for object creation & destruction.
- Per-object map can have problems with duplicate keys (GOOGLE\_unique\_objects layer).

## Object Wrapping

- Complex to implement.
- Modifying - layer must intercept all entry points to unwrap.
- Thread-safe - wrapping happens in-place, no locks required.

# Storing Layer Data – Summary

## Map Lookups

- Simple to implement.
- Non-modifying - only requires lookups.
- Thread-unsafe - requires read-write lock for object creation & destruction.
- Per-object map can have problems with duplicate keys (GOOGLE\_unique\_objects layer).

## Object Wrapping

- Complex to implement.
- Modifying - layer must intercept all entry points to unwrap.
- Thread-safe - wrapping happens in-place, no locks required.
- Supports duplicate object handles.

# Conclusion

- Writing a small layer only takes an hour or two.

# Conclusion

- Writing a small layer only takes an hour or two.
- Once you wrap your head around dispatch tables/chains, the hard part is done.



# Conclusion

- Writing a small layer only takes an hour or two.
- Once you wrap your head around dispatch tables/chains, the hard part is done.
- Layers can do all sorts of fun things - full extensibility of the API in your hands.

# Conclusion

- Writing a small layer only takes an hour or two.
- Once you wrap your head around dispatch tables/chains, the hard part is done.
- Layers can do all sorts of fun things - full extensibility of the API in your hands.
- Questions?

# Appendix

- Full article: <https://renderdoc.org/vulkan-layer-guide.html>
- Sample code: [https://github.com/baldurk/sample\\_layer](https://github.com/baldurk/sample_layer)

Thanks to Matthäus Chajdas, Jasper Bekkers for checking slides :).